

Efficient Query Processing for Spatial and Temporal Data Exploration

Eleni Tzirita Zacharatou

Advisor: Anastasia Ailamaki

Public Thesis Defense

09.08.2019

Urbanization and City Planning



Data Exhaust from Cities

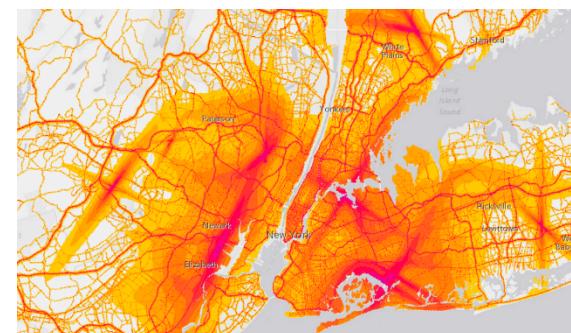
Infrastructure



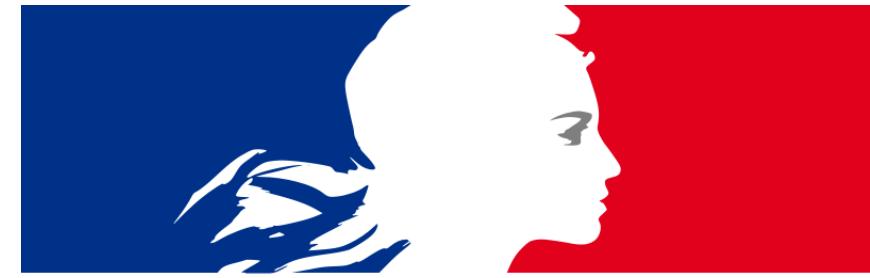
Environment



People

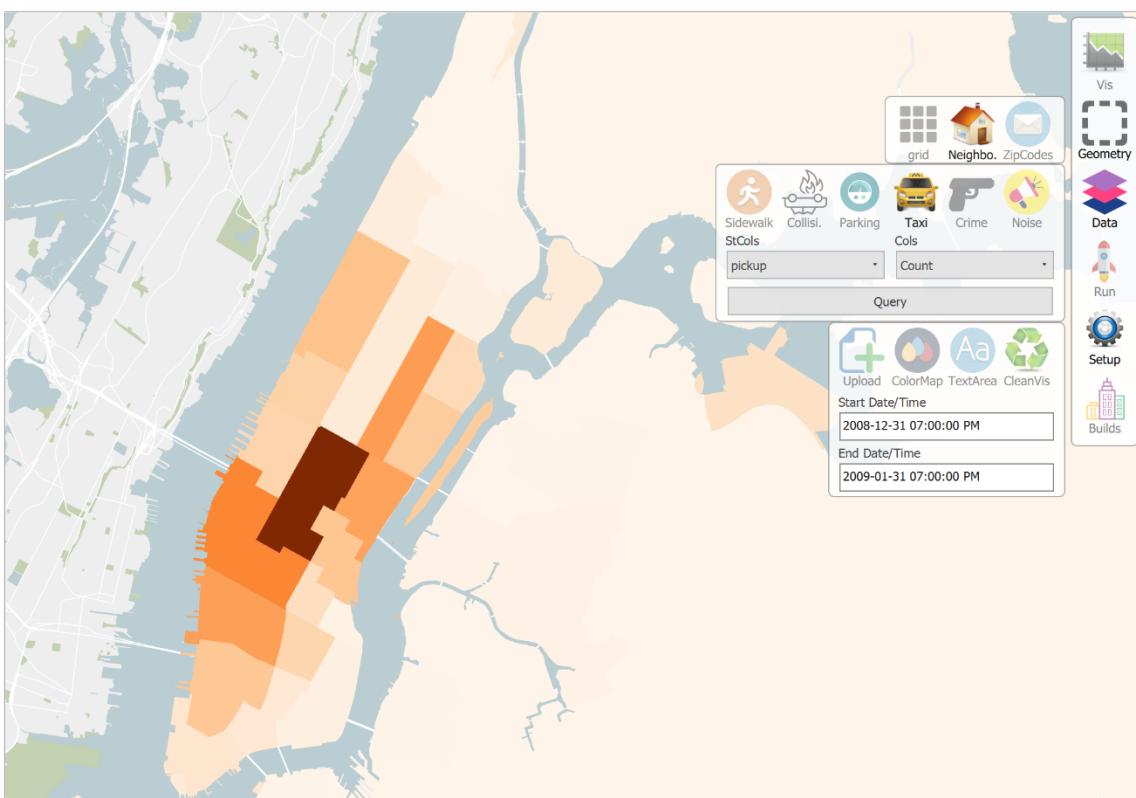


Understanding Cities through Data

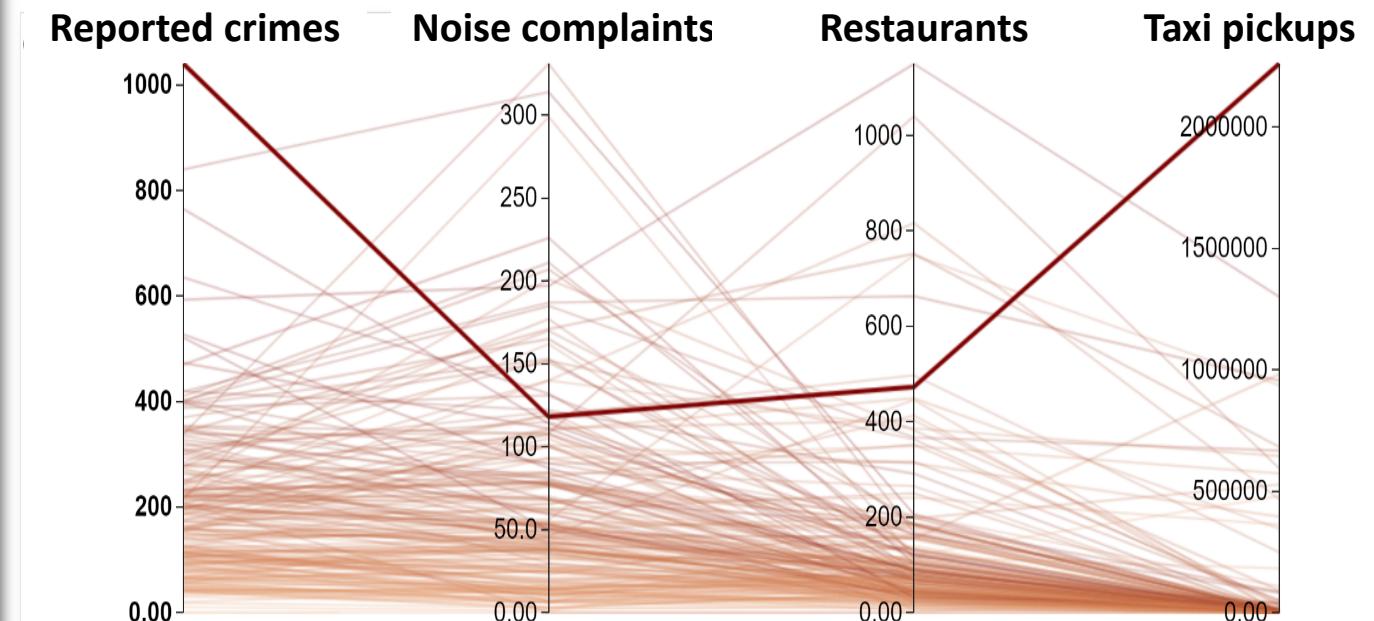


Opportunity: Data-driven urban planning

Visual Spatial Data Exploration



Distribution of taxi pickups per neighborhood in Manhattan



Comparison of different urban datasets

Need: Interactive response times

Spatial Aggregation Queries

Aggregation

SELECT COUNT(*)

Input

FROM taxi ride T , neighborhoods N

Spatial Join

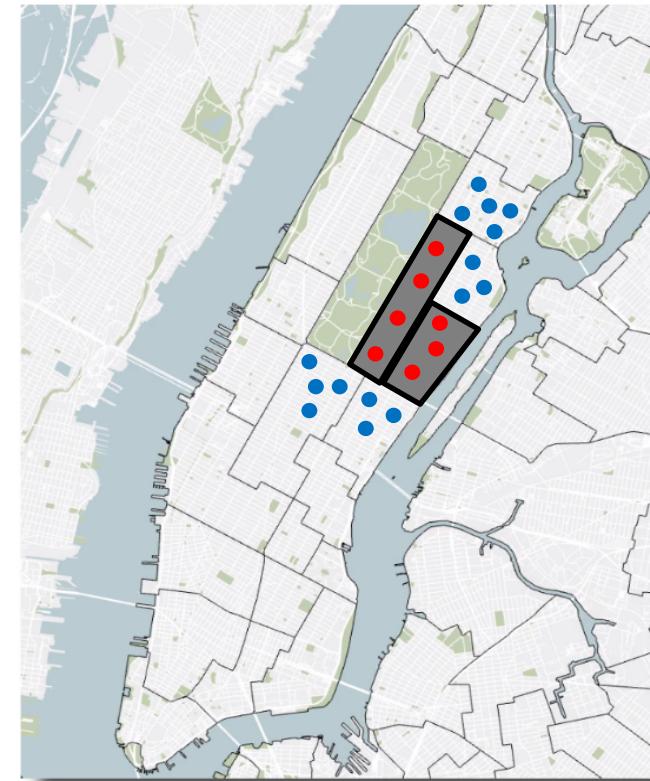
WHERE $T.\text{pickup}$ INSIDE $N.\text{geometry}$

Selection

AND $T.\text{picktime}$ in January 2009

Grouping

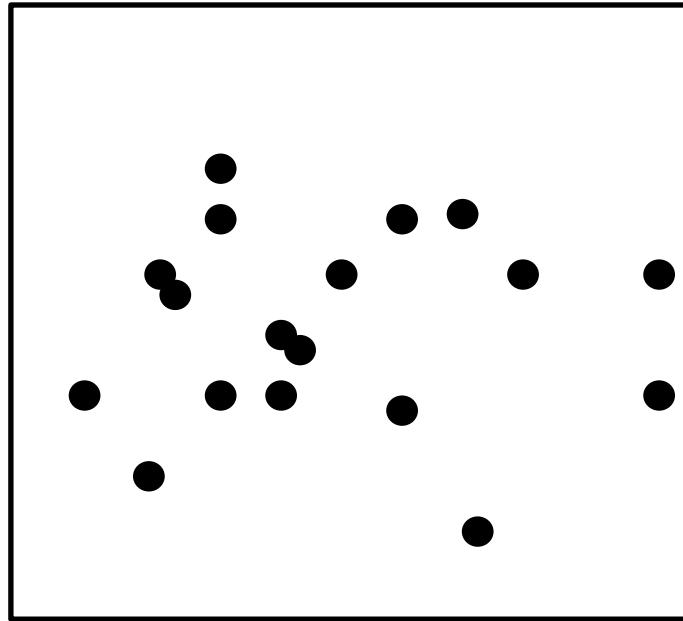
GROUP BY $N.\text{id}$



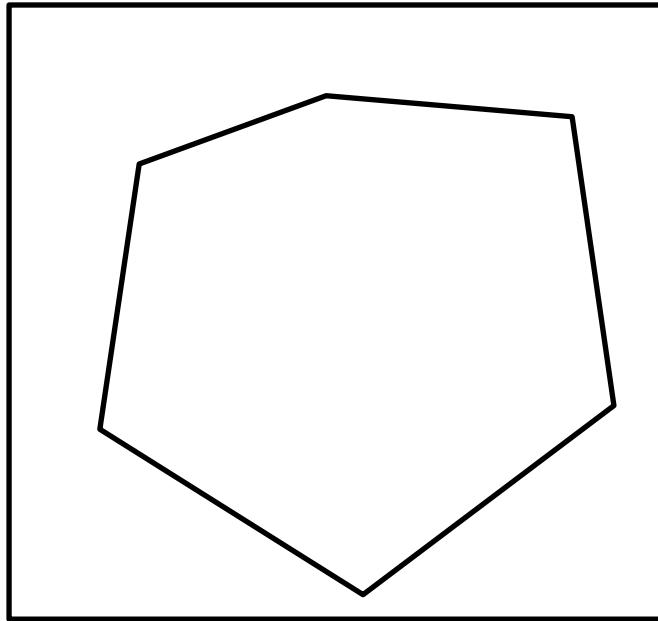
Point-in-Polygon tests

Expensive Point-in-Polygon tests → High latency (minutes)

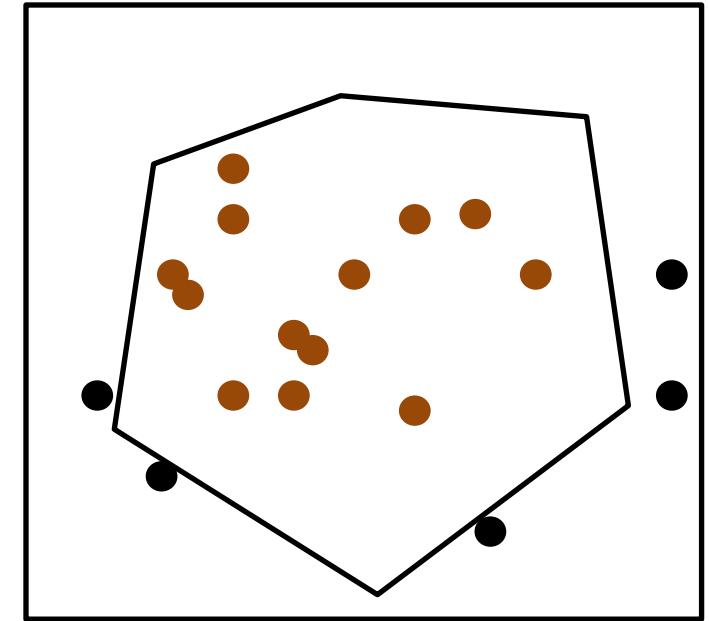
Spatial Aggregation: a Geometric Perspective



Input points



Input polygon

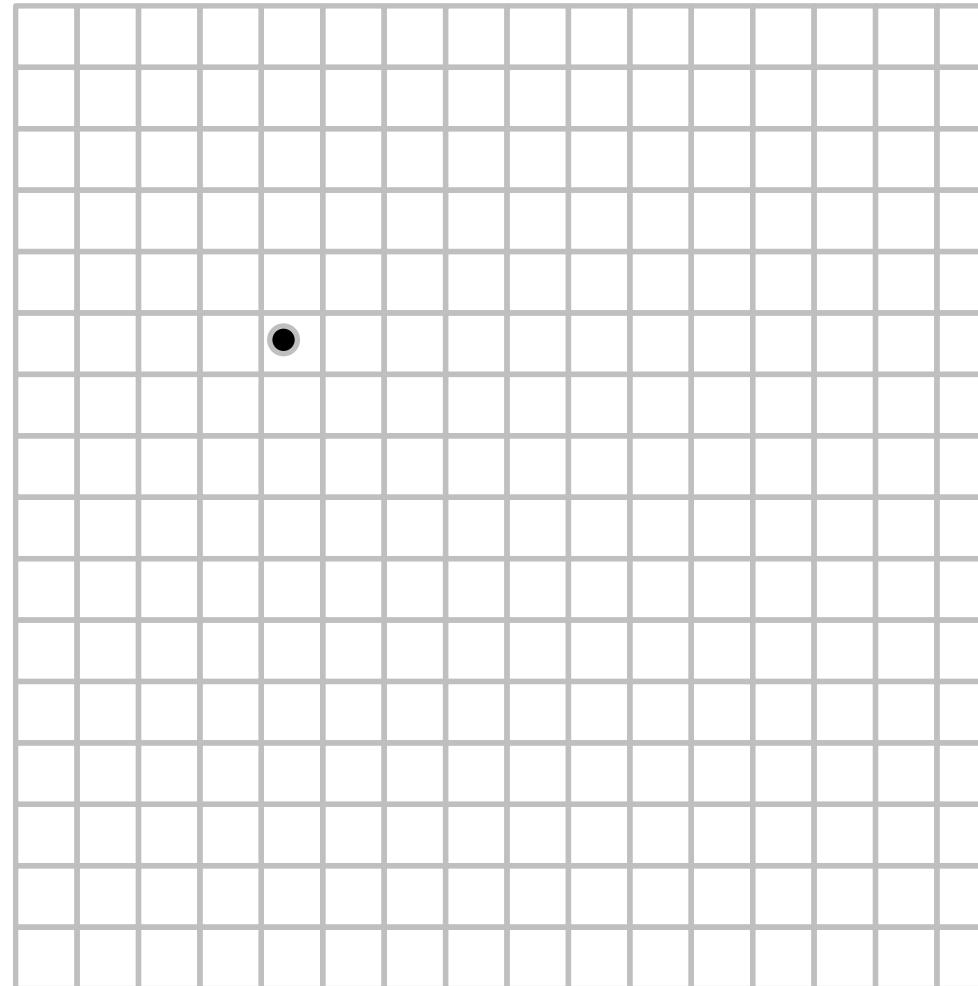


Spatial join

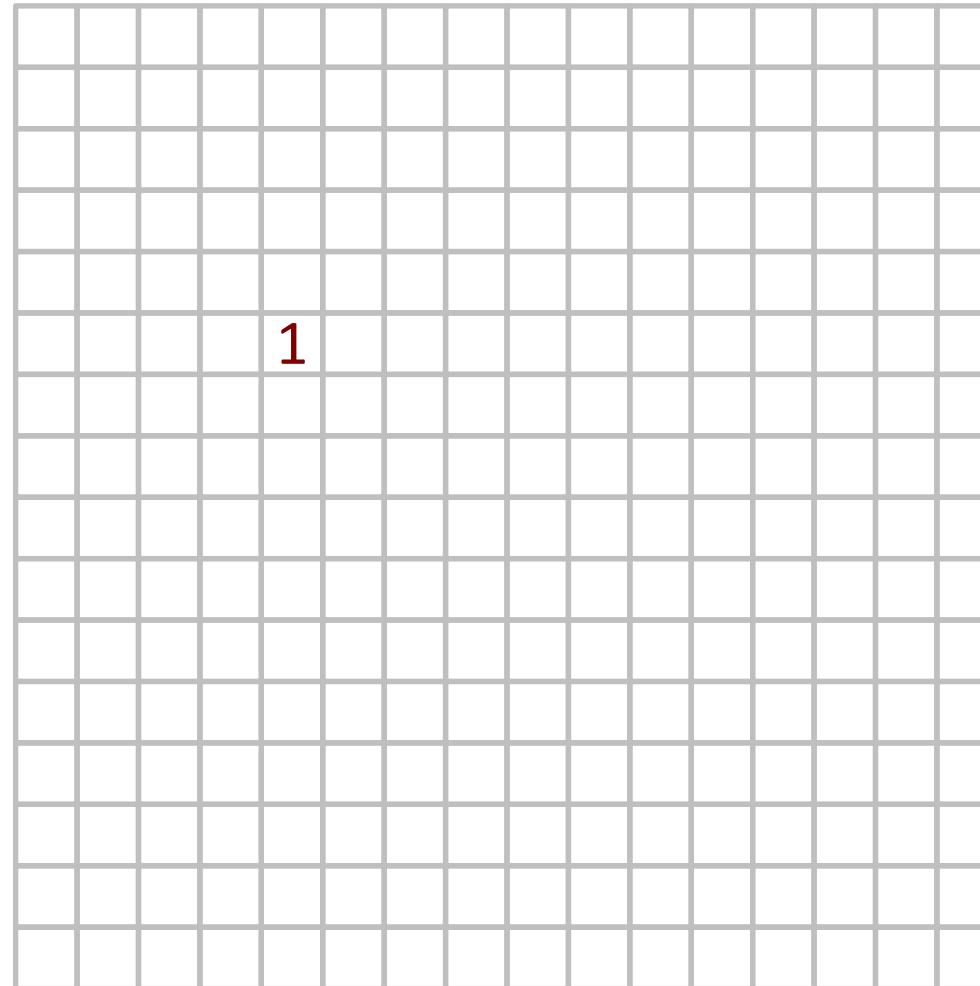
“Drawing” on the same canvas

→ Leverage the graphics pipeline of the GPU

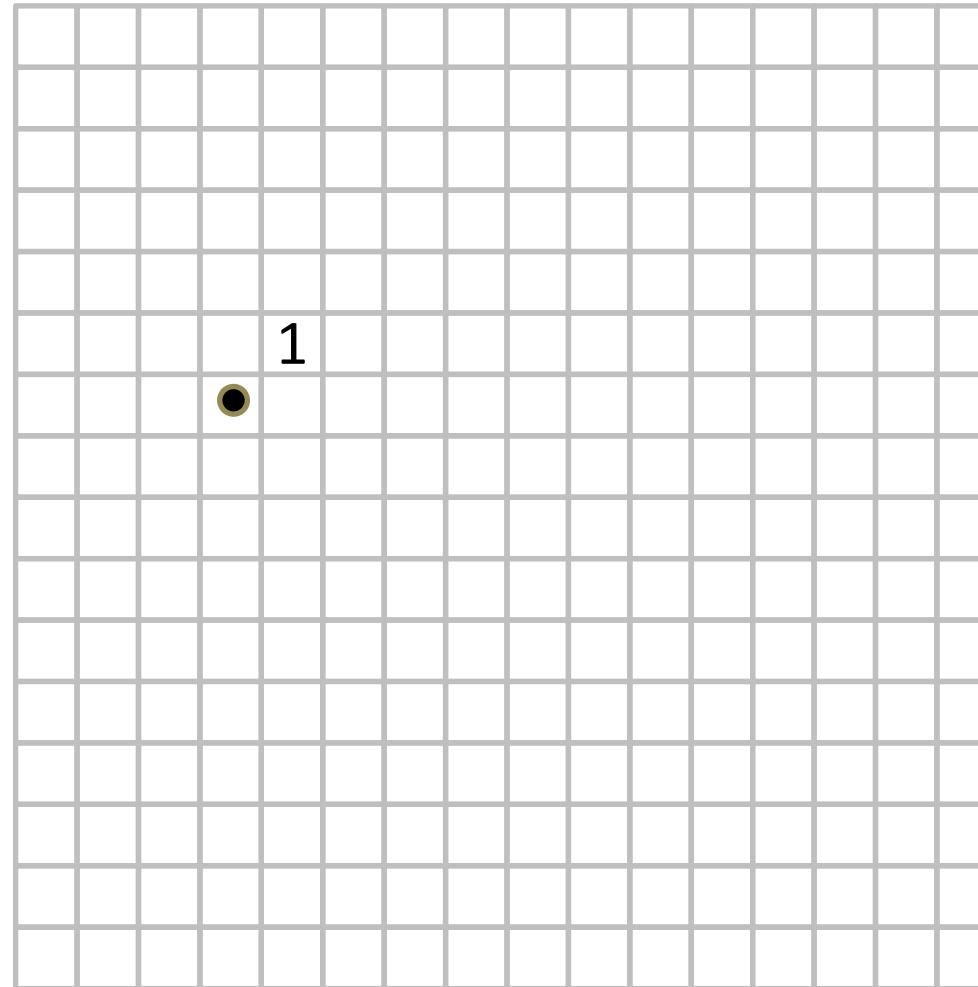
Raster Join - Step 1: Draw the Points



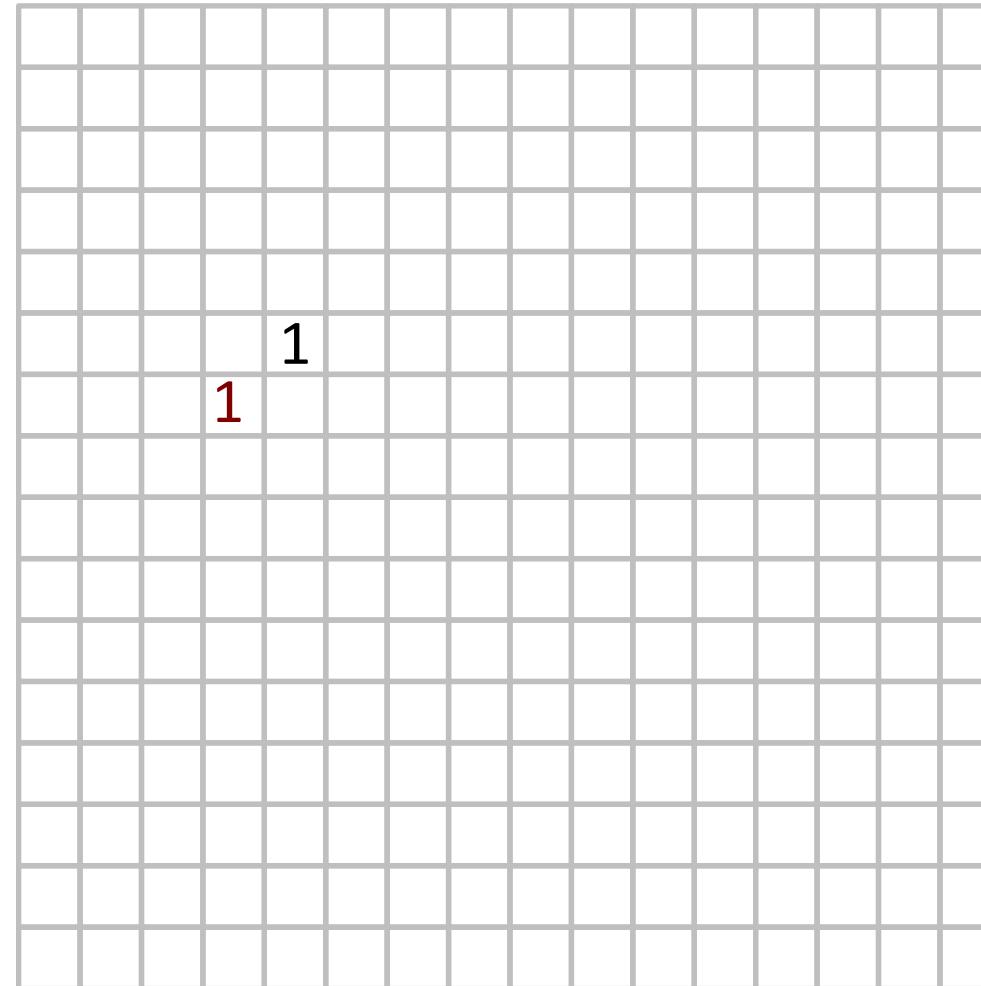
Raster Join - Step 1: Draw the Points



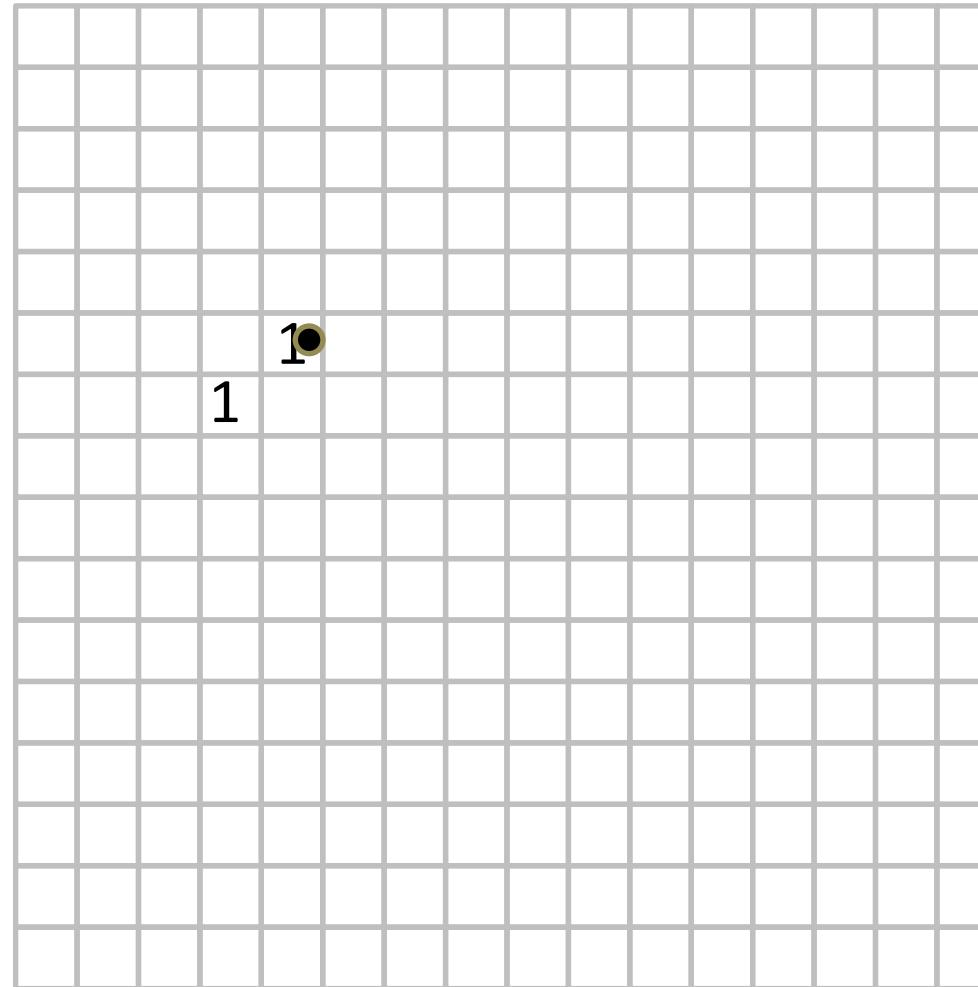
Raster Join - Step 1: Draw the Points



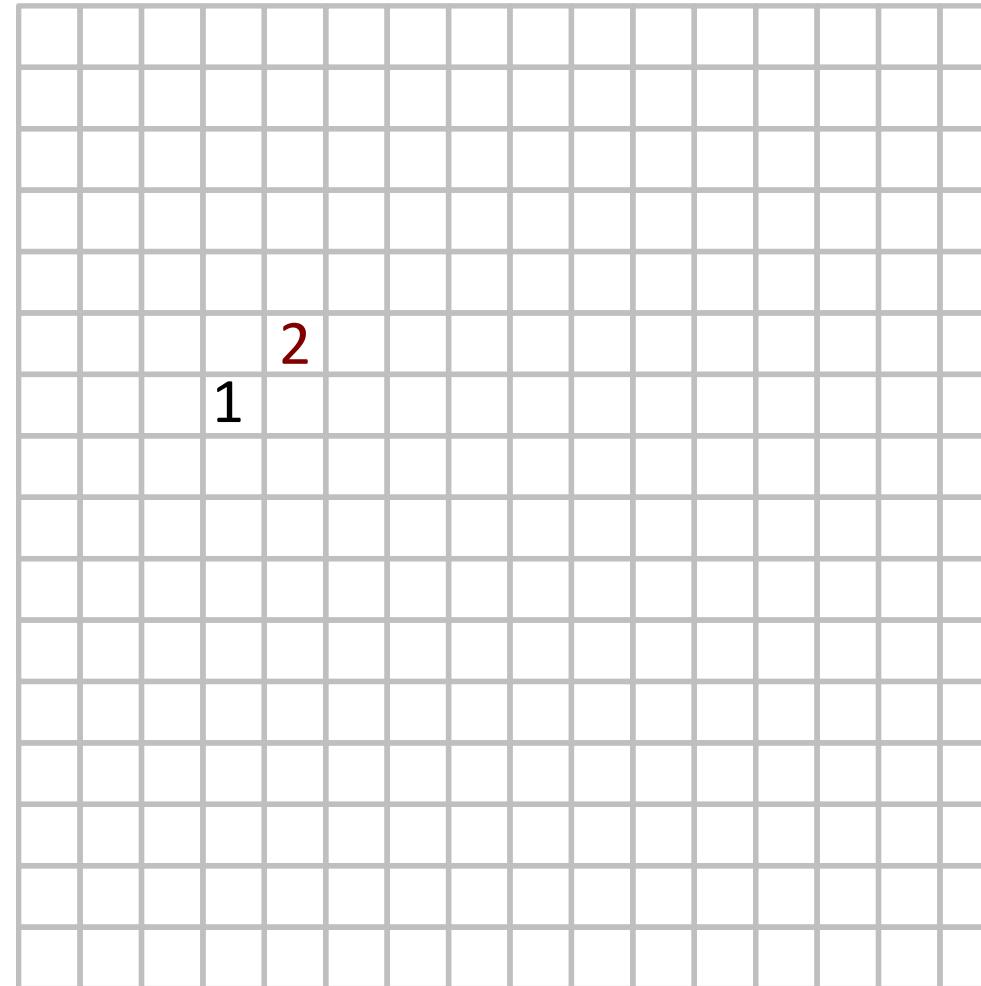
Raster Join - Step 1: Draw the Points



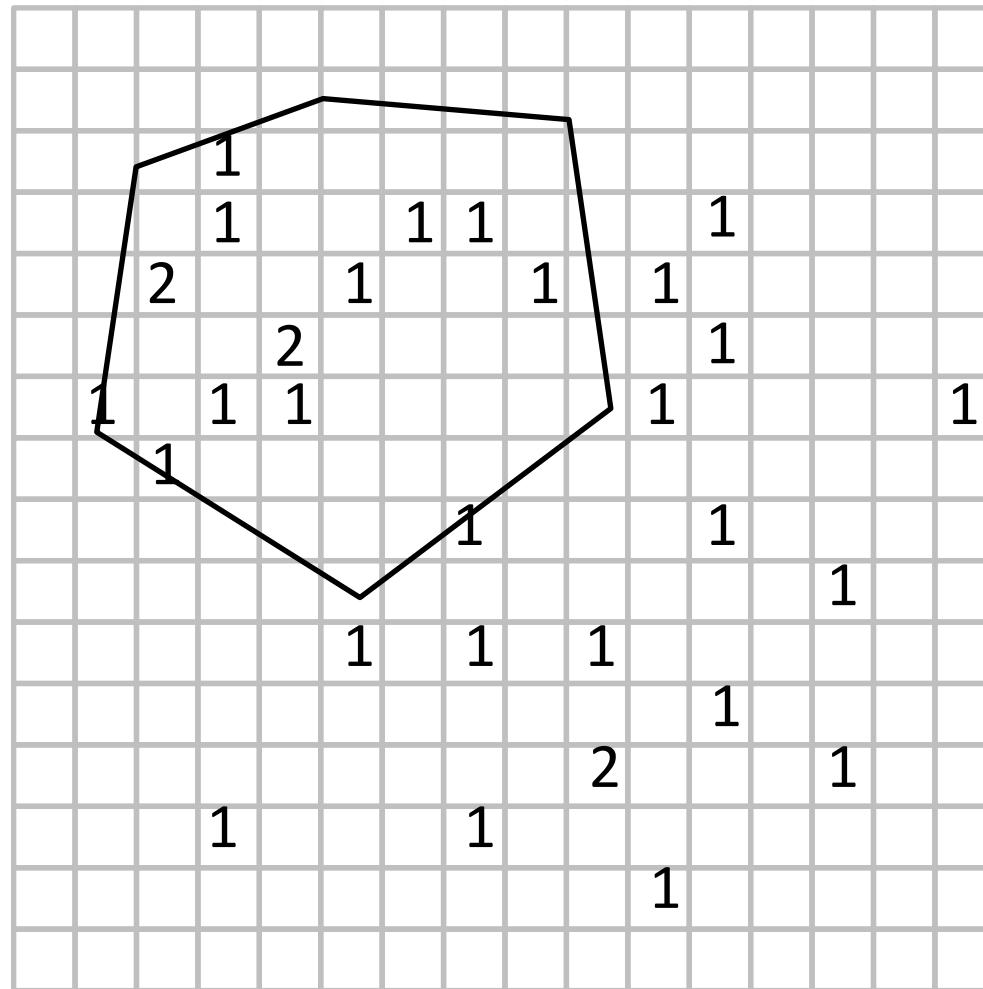
Raster Join - Step 1: Draw the Points



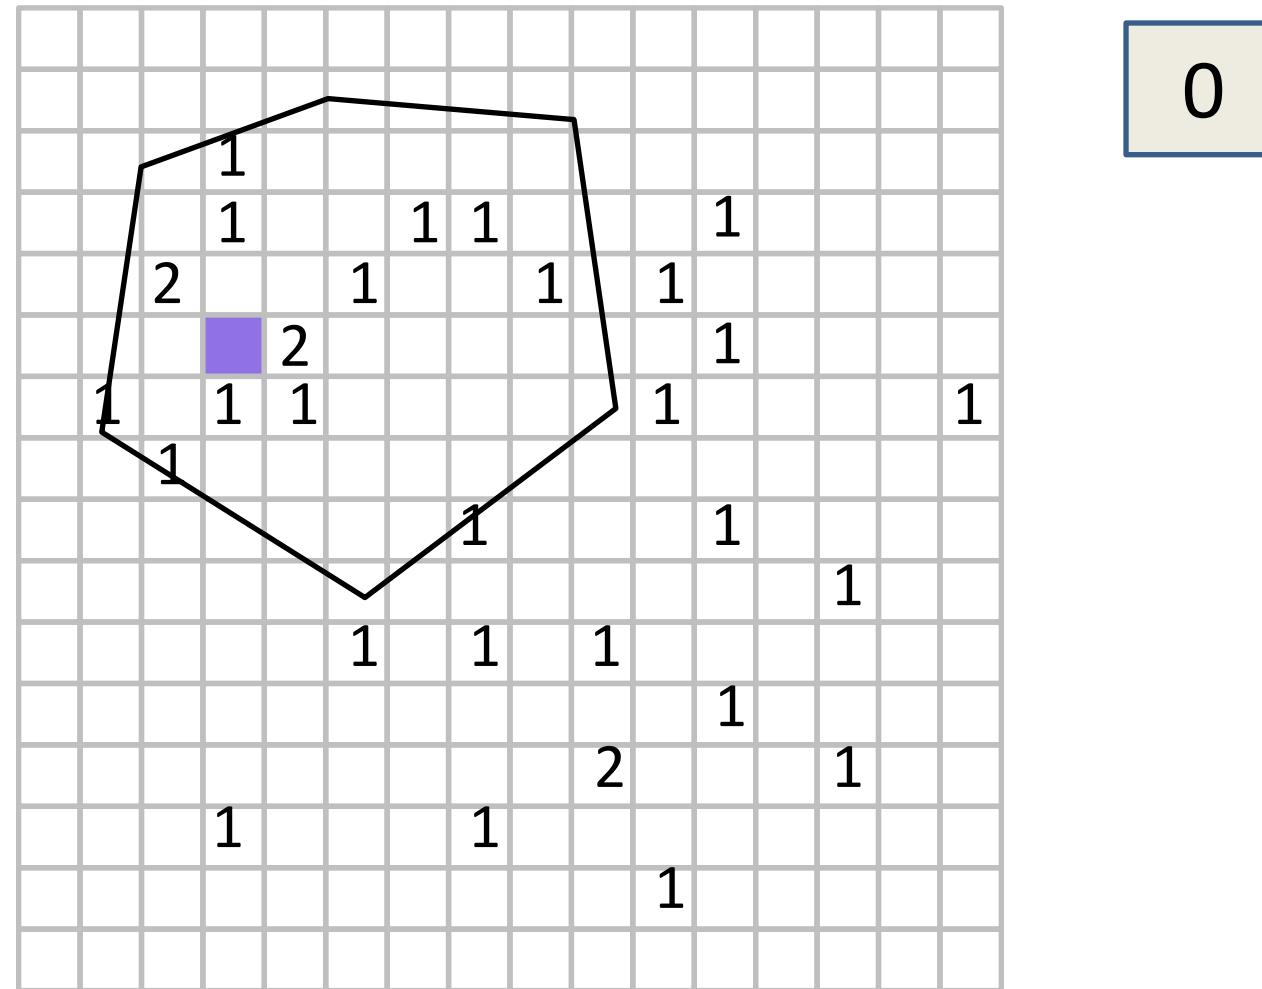
Raster Join - Step 1: Draw the Points



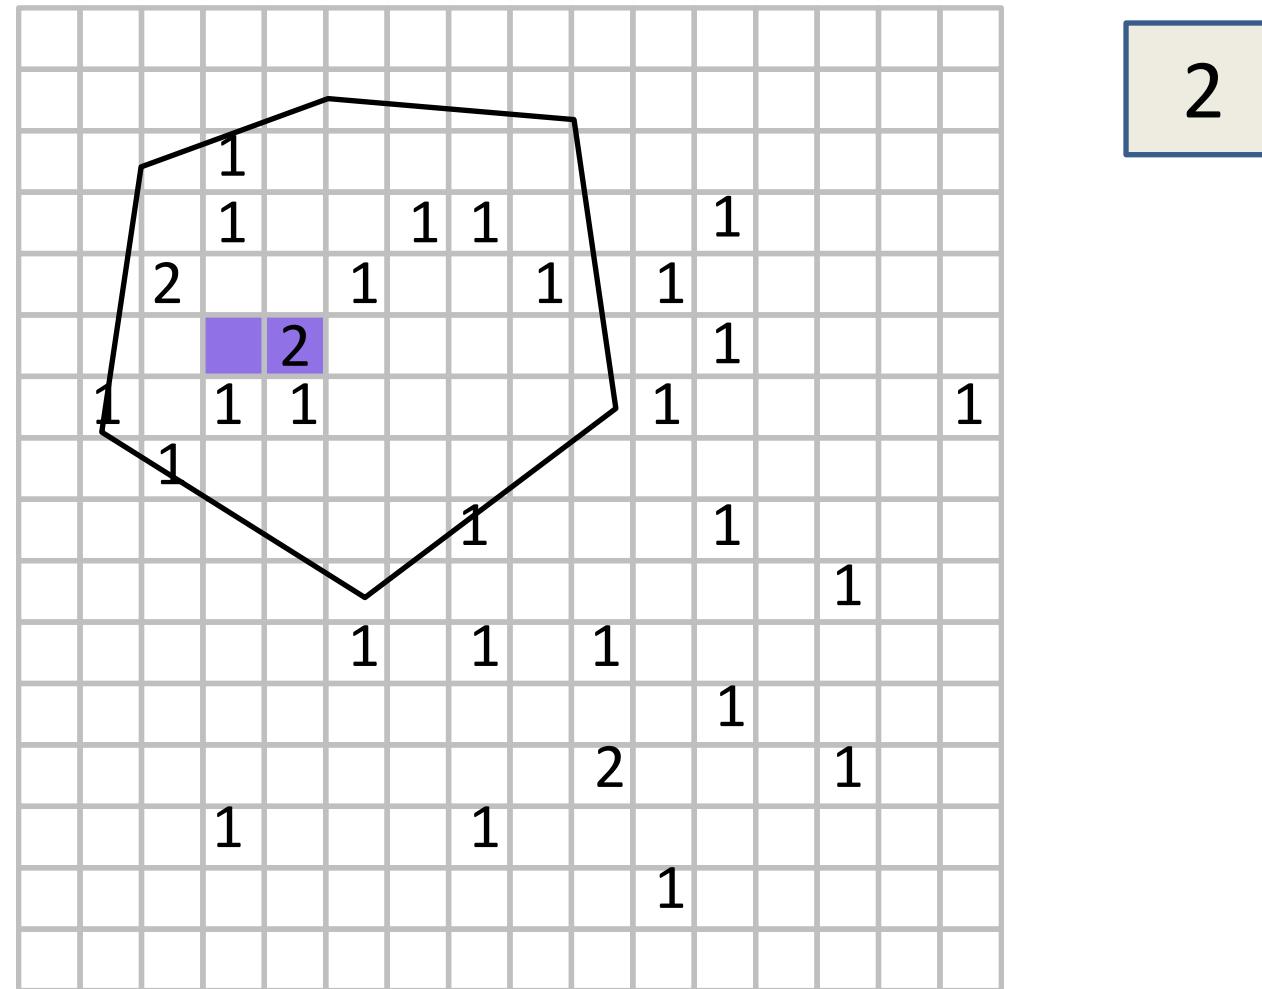
Raster Join - Step 1: Draw the Points



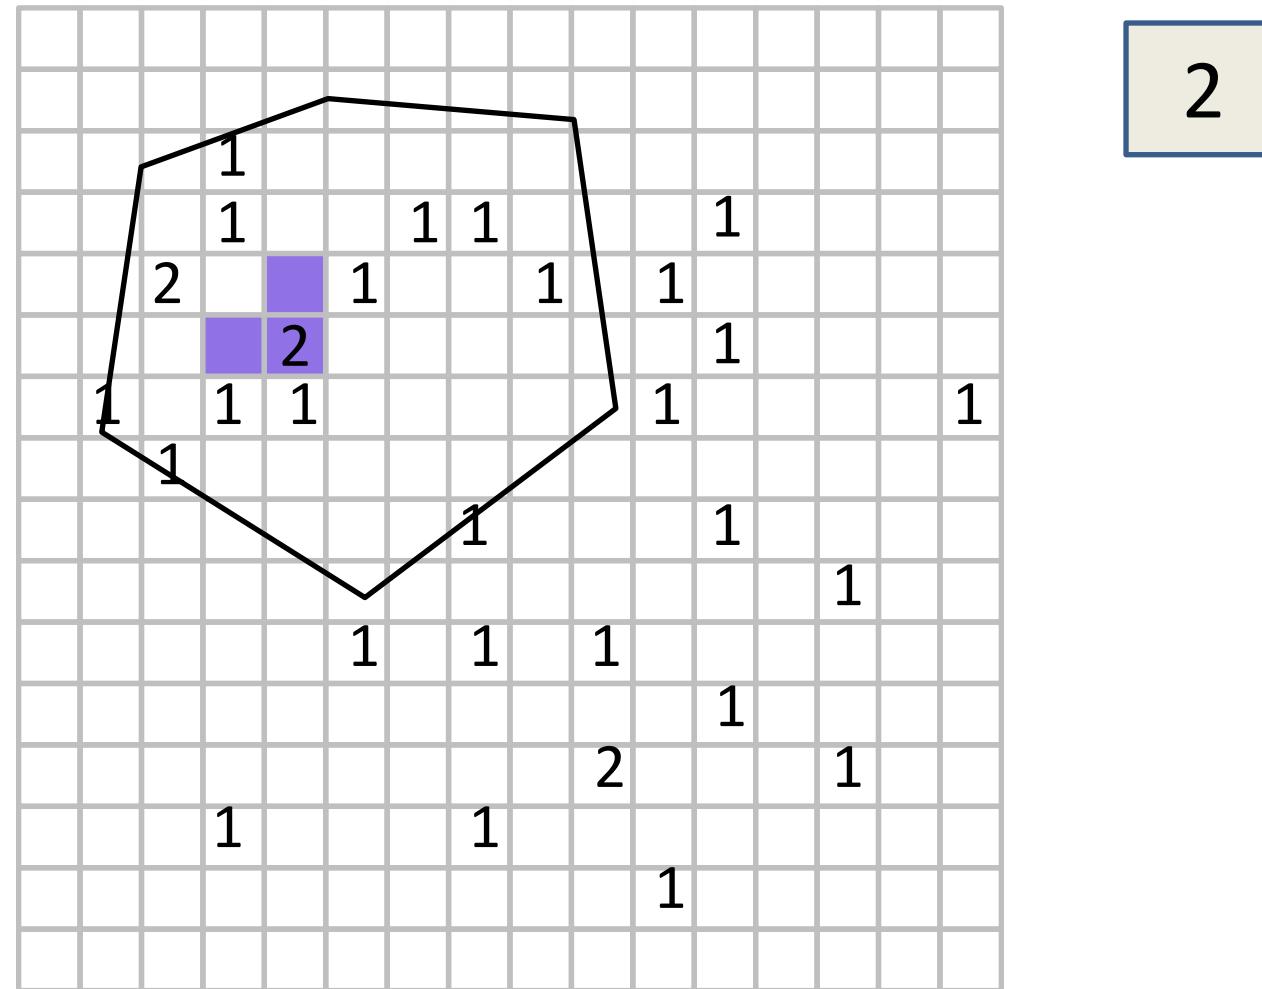
Raster Join - Step 2: Draw the Polygons



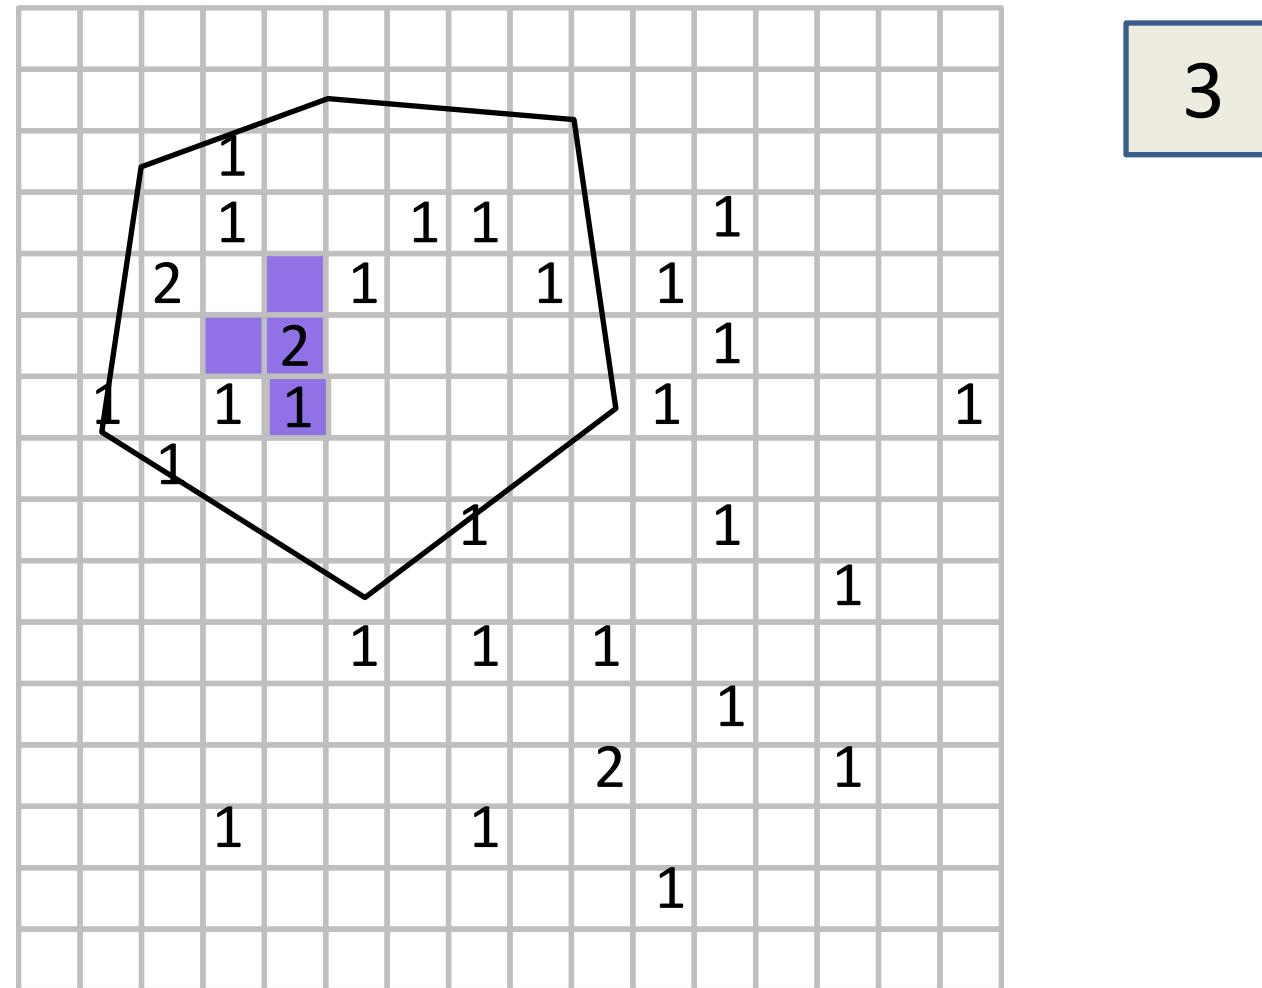
Raster Join - Step 2: Draw the Polygons



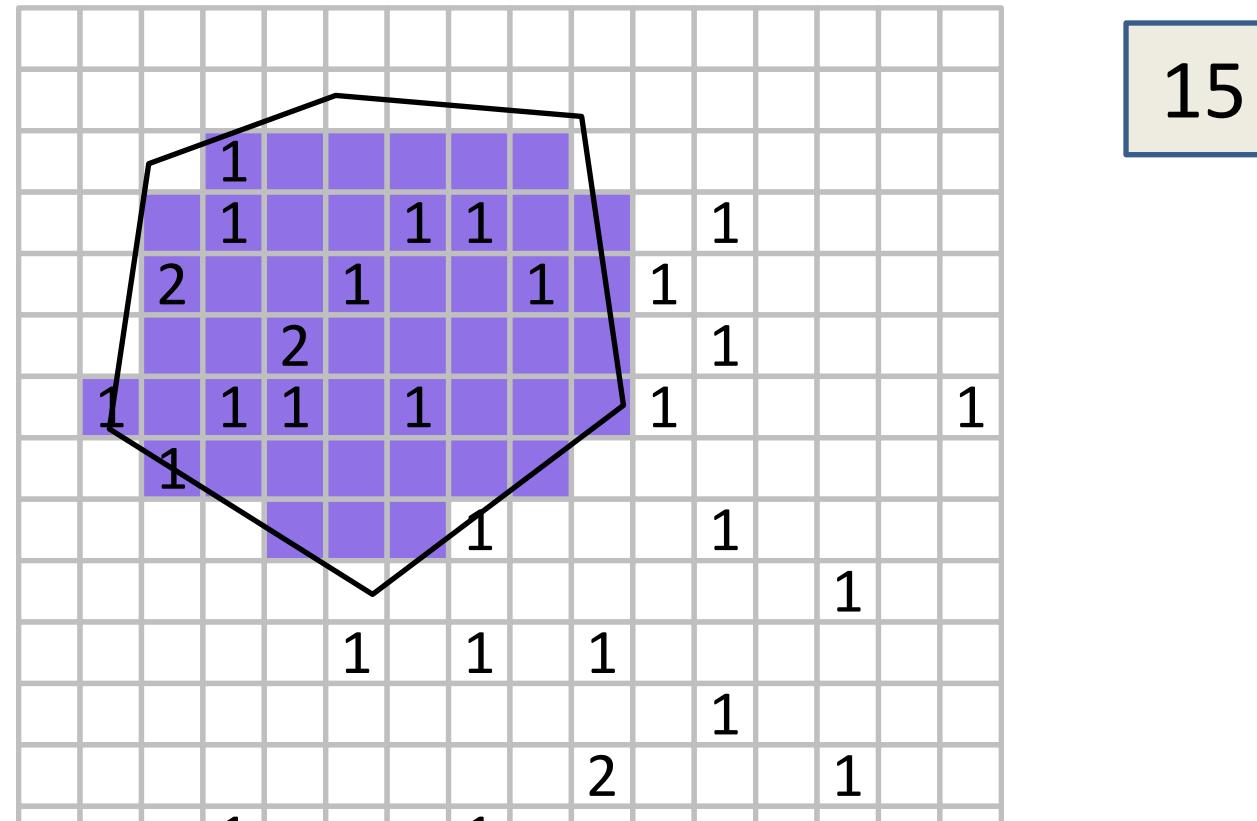
Raster Join - Step 2: Draw the Polygons



Raster Join - Step 2: Draw the Polygons



Raster Join - Step 2: Draw the Polygons



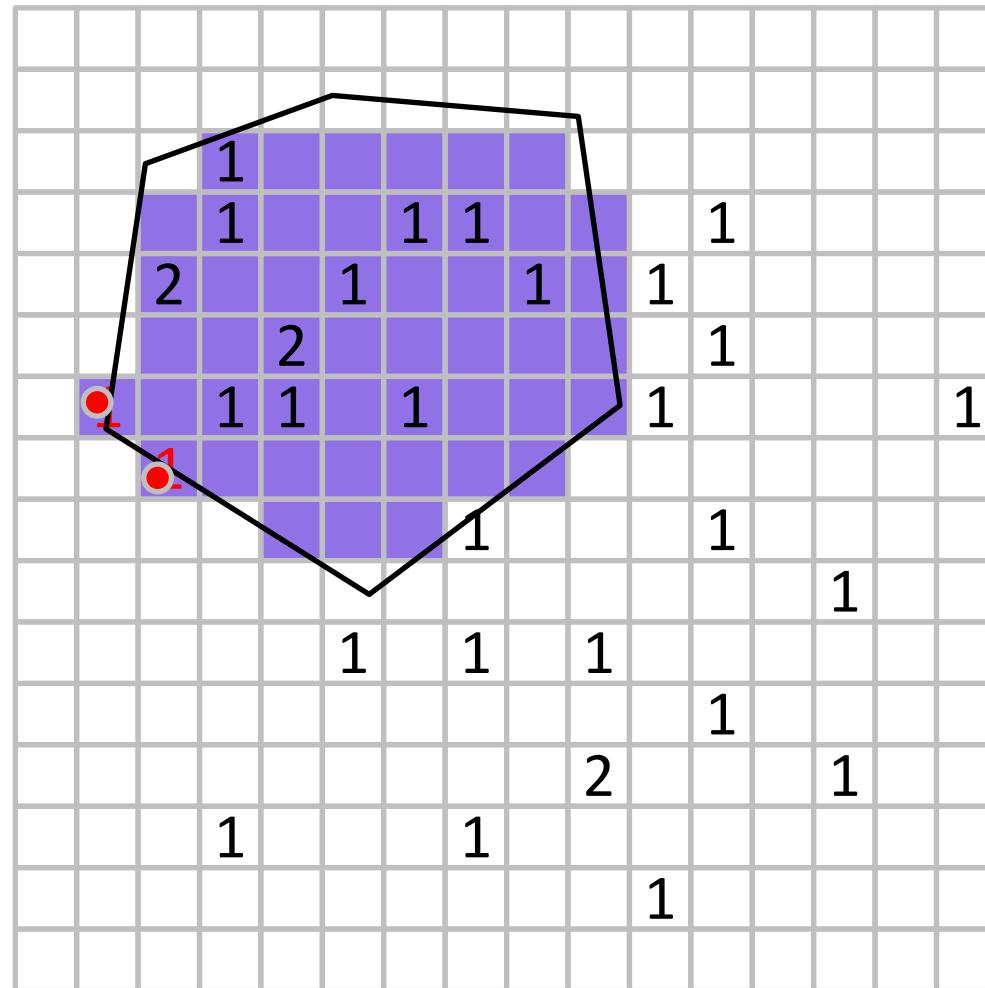
Uses native support for drawing in GPUs
Combines the aggregation with the join operation
No Point-in-Polygon tests

Bounding the Approximation Error

- Bound the Hausdorff distance between the approximate (purple) and the original polygon.

$$H(P_a, P) \leq \varepsilon$$

- Smaller pixel size → higher accuracy.

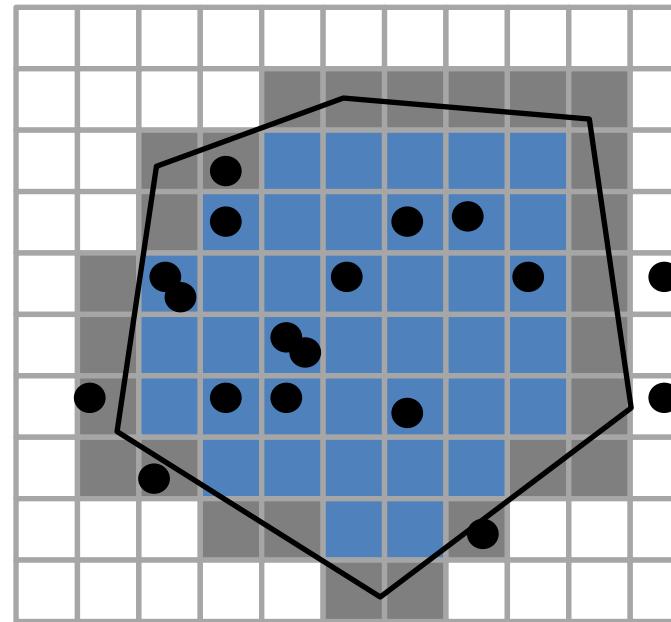


Trade accuracy for response time

Hybrid Raster Join: an Accurate Technique

Blue pixels - completely inside the polygon: store count

Grey pixels - polygon boundary: Point-in-Polygon (PiP) tests



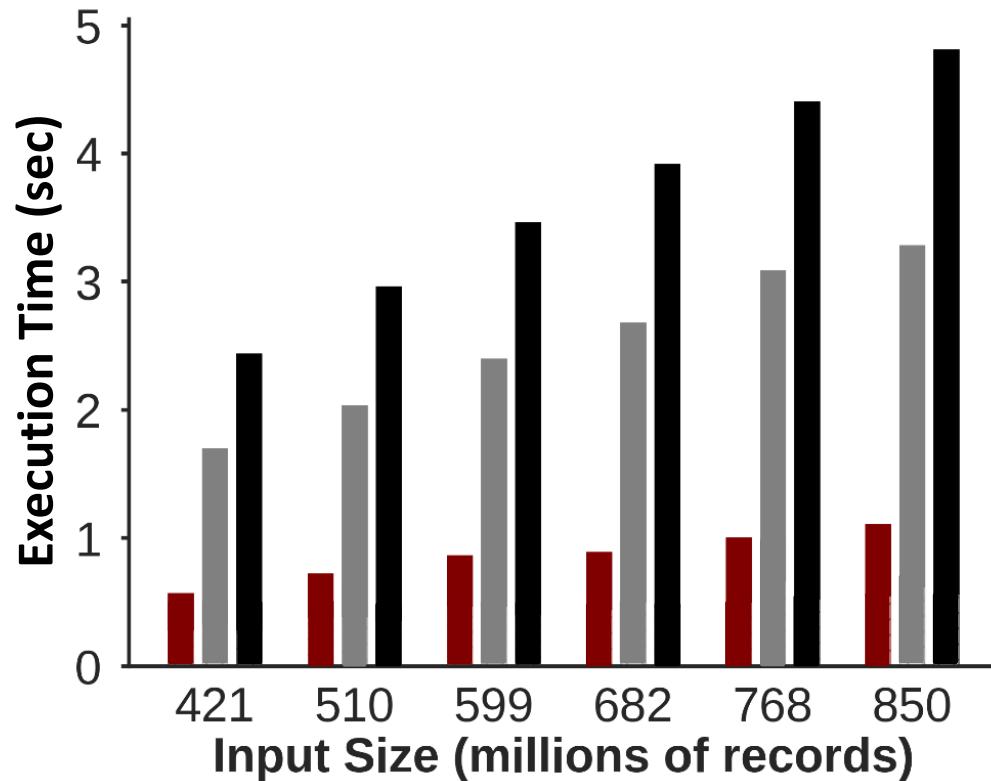
Extra computation: identifying the boundary & performing PiP tests

Scaling with Increasing Data Sizes

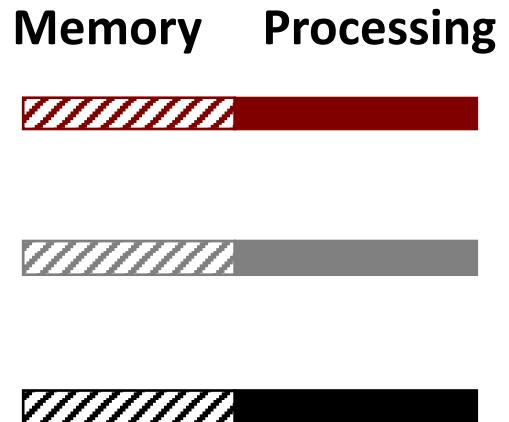
COUNT Taxi rides (points) GROUP BY NYC Neighborhoods (260 polygons)

[Intel Core i7 Quad-Core CPU @ 2.80GHz, 16GB RAM, NVIDIA GTX 1060 GPU, 6GB memory (using only 3GB)]

[Max resolution: 8192, Grid index: 1024^2 cells]



- Bounded Raster Join**
(Approximate, ϵ -bound = 10 meters)
- Hybrid Raster Join**
(Accurate)
- GPU Baseline**
(Accurate)

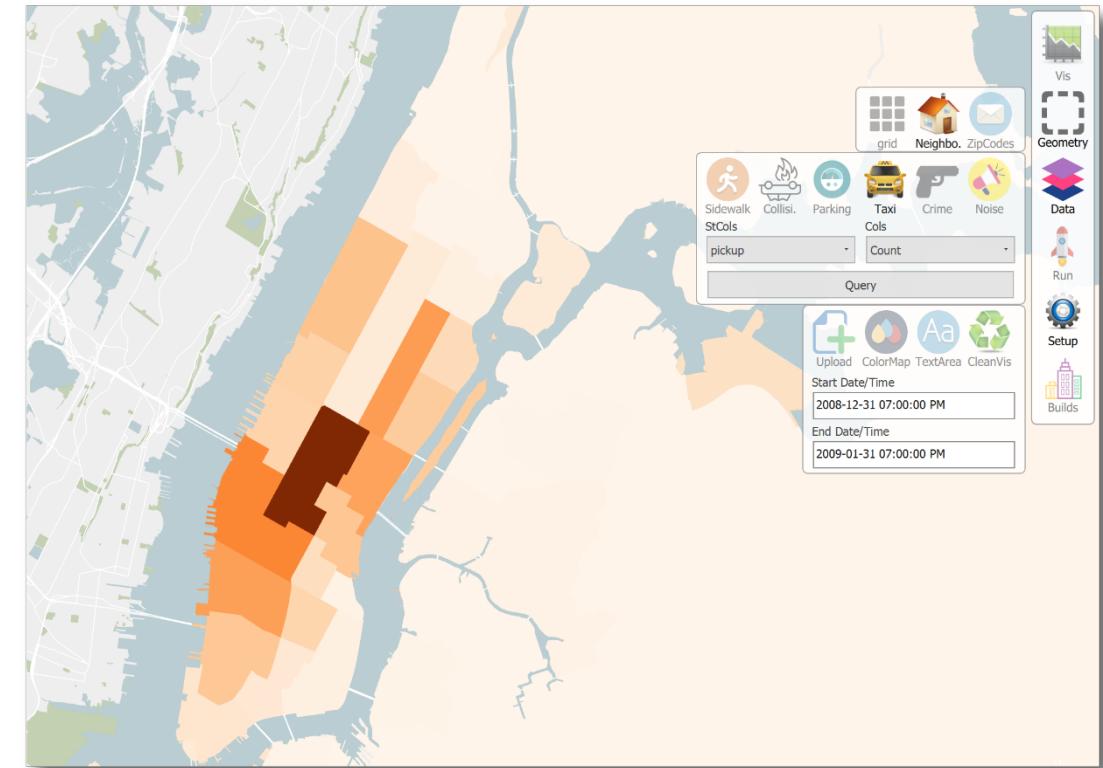


Bounded Raster Join is 4X faster than GPU Baseline
CPU-GPU data transfer takes a significant amount of time

GPU Rasterization enables Interactive Spatial Queries

[VLDB18, SIGMOD18]

- Express queries as graphics primitives and use modern GPUs
- Aggregating 850M taxi records over NYC neighborhoods in ~1 second

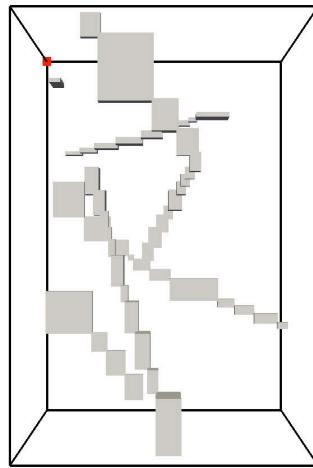


Clipped Minimum Bounding Boxes for Efficient Spatial Indexing

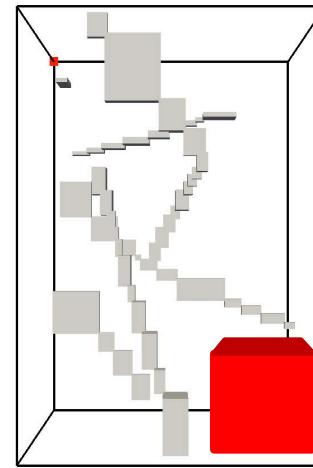
[ICDE18]

Improve precision by subtracting out empty bounded areas

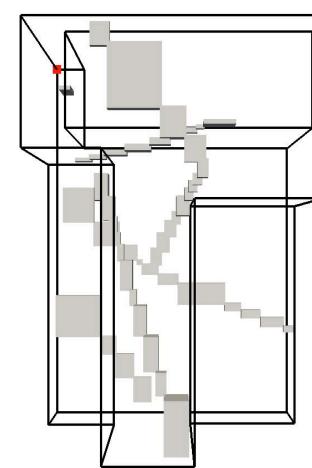
→ Answering a spatial range query on 1B objects in less than 200ms



Minimum Bounding Box



Empty space →
Ineffective filtering



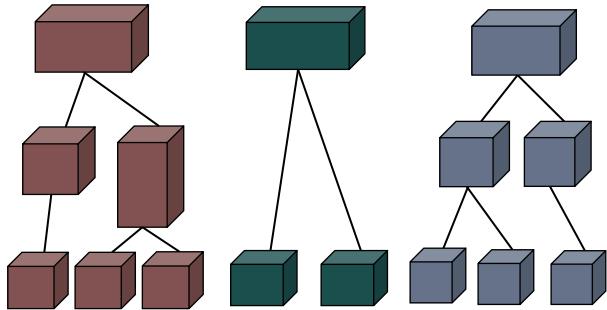
Clipped Bounding Box

Workload-Aware Indexing enables Ad-hoc Spatial Queries

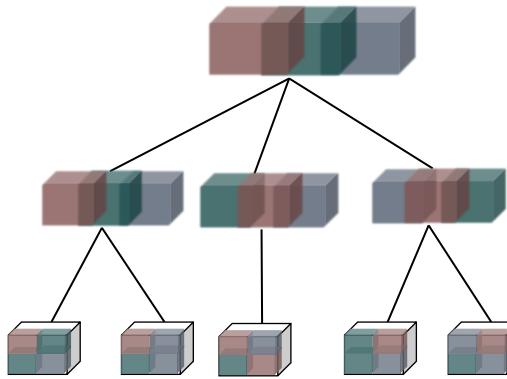
[ExploreDB16]

Category-aware spatial data organization

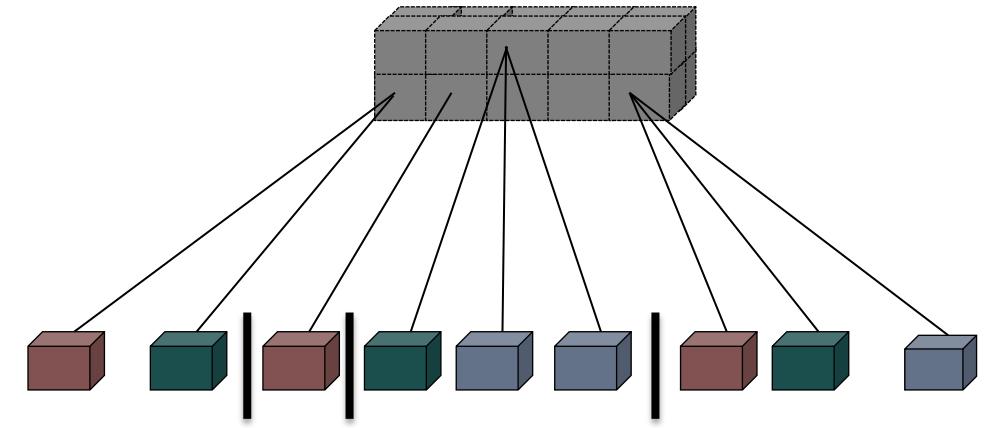
→ Up to 12.3X faster queries on 10 different neuron categories



Index per category



Index over union



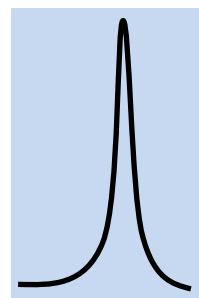
Distinct group per category

Quadtree Bitmap Decomposition for Scalable Time Series Indexing

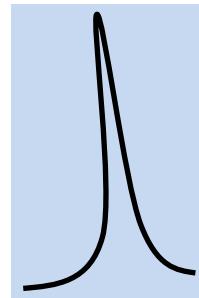
Value-time searches exploiting space-time similarity

[SSDBM15]

→ From 9X to 23X faster queries on neuroscience data



0	0	0	0
0	0	1	0
1	1	1	1
1	1	1	1



0	0	0	0
0	1	0	0
1	1	1	1
1	1	1	1

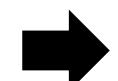
Bitmap
encoding



0	0	0	0
0	0	1	0
1	1	1	1
1	1	1	1

0	0	0	0
0	1	0	0
1	1	1	1
1	1	1	1

Quadtree decomposition
Exploit intra-similarities



0	0	0	0
0	0	1	0
1	1	1	1
1	1	1	1

0	0	0	0
0	1	0	0
1	1	1	1
1	1	1	1

Bitmap grouping
Exploit inter-similarities

Thesis Statement

Modern applications need to explore large amounts of spatial and temporal data at interactive speeds, challenging traditional query processing techniques that rely on time-consuming computations and inefficient access methods.

Query operators that exploit **specialized hardware** and **workload-aware access methods** enable scalable and interactive exploration of spatial and temporal data.

Looking Ahead

- **Approximation-based spatial data processing**
 - Fine-grained approximations and omission of exact geometric tests
 - Distance-based error bound
 - Trade precision / storage space for performance
- **Utility of graphics techniques for spatial data processing**
 - GPU rasterization for real-time approximation
 - 3D Join → Collision Detection

Thank you!